

# TidyVerse

Zachary Safir

4/10/2021

```
library(tidyverse)
library(knitr)
library(kableExtra)
```

## Introduction

The tidyverse contains a collection of data science packages that work together in harmony to accomplish various goals. This vignette will demonstrate several ways to make full use of their combined capability.

## The Data

For this demonstration, we will use a dataset that is included with dplyr itself. It contains data on the characters from the Starwars series. It contains various pieces of data that describe each character.

```
kable(head(starwars),format = "latex") %>%
  kable_styling(wraptable_width = "0pt")
```

starwars

```
## # A tibble: 87 x 14
##   name      height mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Luke S~    172    77 blond      fair        blue        19   male masculin
## 2 C-3PO      167    75 <NA>      gold        yellow      112  none masculin
## 3 R2-D2       96    32 <NA>      white, bl~ red        33   none masculin
## 4 Darth ~    202   136 none      white       yellow     41.9 male masculin
## 5 Leia O~    150    49 brown     light       brown       19   fema~ feminin
## 6 Owen L~    178   120 brown, grey light       blue       52   male masculin
## 7 Beru W~    165    75 brown     light       blue       47   fema~ feminin
## 8 R5-D4       97    32 <NA>      white, red red        NA   none masculin
## 9 Biggs ~    183    84 black     light       brown       24   male masculin
## 10 Obi-Wa~   182    77 auburn, wh~ fair        blue-gray   57   male masculin
## # ... with 77 more rows, and 5 more variables: homeworld <chr>, species <chr>,
## #   films <list>, vehicles <list>, starships <list>
```

Interestingly, some of the columns of data are full of lists. The column displayed below, shows which films a character appeared in.

```
head(starwars$films)
```

```
## [[1]]
## [1] "The Empire Strikes Back" "Revenge of the Sith"
## [3] "Return of the Jedi"      "A New Hope"
## [5] "The Force Awakens"
##
## [[2]]
## [1] "The Empire Strikes Back" "Attack of the Clones"
## [3] "The Phantom Menace"      "Revenge of the Sith"
## [5] "Return of the Jedi"      "A New Hope"
##
## [[3]]
## [1] "The Empire Strikes Back" "Attack of the Clones"
## [3] "The Phantom Menace"      "Revenge of the Sith"
## [5] "Return of the Jedi"      "A New Hope"
## [7] "The Force Awakens"
##
## [[4]]
## [1] "The Empire Strikes Back" "Revenge of the Sith"
## [3] "Return of the Jedi"      "A New Hope"
##
## [[5]]
## [1] "The Empire Strikes Back" "Revenge of the Sith"
## [3] "Return of the Jedi"      "A New Hope"
## [5] "The Force Awakens"
##
## [[6]]
## [1] "Attack of the Clones" "Revenge of the Sith" "A New Hope"
```

The first thing to figure out is how to pick out only characters that appear in certain films. In order to use filter from dplyr on a list, we need to use a purr function with it. As filter is expecting a logical value, we need to return something logical. Using map\_lgl, we can accomplish this.

```
starwars %>%
  filter(map_lgl(films, ~ "Attack of the Clones" %in% .))
```

```
## # A tibble: 40 x 14
##   name      height mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>    <int> <dbl> <chr>    <chr>    <chr>    <dbl> <chr> <chr>
## 1 C-3PO      167  75  <NA>    gold     yellow     112  none  mascu~
## 2 R2-D2       96  32  <NA>    white, bl~ red        33  none  mascu~
## 3 Owen L~    178 120  brown, grey light     blue       52  male  mascu~
## 4 Beru W~    165  75  brown    light     blue       47  fema~  femin~
## 5 Obi-Wa~    182  77  auburn, wh~ fair      blue-gray   57  male  mascu~
## 6 Anakin~    188  84  blond    fair      blue      41.9 male  mascu~
## 7 Yoda       66  17  white    green     brown     896  male  mascu~
## 8 Palpat~    170  75  grey     pale      yellow     82  male  mascu~
## 9 Boba F~    183 78.2  black    fair      brown     31.5 male  mascu~
## 10 Nute G~    191  90  none     mottled g~ red       NA  male  mascu~
```

```
## # ... with 30 more rows, and 5 more variables: homeworld <chr>, species <chr>,
## #   films <list>, vehicles <list>, starships <list>
```

In order to use filter on multiple values, we need to use the base R function “all”.

```
starwars %>%
  filter(map_lgl(films, ~ all( c("Attack of the Clones", "A New Hope") %in% .)))

## # A tibble: 5 x 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 C-3PO      167    75 <NA>      gold       yellow      112 none masculin
## 2 R2-D2      96    32 <NA>      white, bl~ red        33 none masculin
## 3 Owen La~   178   120 brown, grey light      blue        52 male masculin
## 4 Beru Wh~   165    75 brown      light      blue        47 fema~ feminin
## 5 Obi-Wan~   182    77 auburn, wh~ fair       blue-gray   57 male masculin
## # ... with 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

We can also use tidyr in order to flatten our lists full of data out. The resulting dataframe of this action is shown below.

```
starwars %>%
  select(name, films) %>%
  unnest(films)

## # A tibble: 173 x 2
##   name      films
##   <chr>      <chr>
## 1 Luke Skywalker The Empire Strikes Back
## 2 Luke Skywalker Revenge of the Sith
## 3 Luke Skywalker Return of the Jedi
## 4 Luke Skywalker A New Hope
## 5 Luke Skywalker The Force Awakens
## 6 C-3PO      The Empire Strikes Back
## 7 C-3PO      Attack of the Clones
## 8 C-3PO      The Phantom Menace
## 9 C-3PO      Revenge of the Sith
## 10 C-3PO     Return of the Jedi
## # ... with 163 more rows
```

With our data in a normal format, we can use the dplyr count function to discover which film is most common.

```
starwars %>%
  unnest(films) %>%
  count(films) %>%
  arrange(n)
```

```
## # A tibble: 7 x 2
##   films          n
##   <chr>        <int>
## 1 The Force Awakens    11
## 2 The Empire Strikes Back 16
## 3 A New Hope          18
## 4 Return of the Jedi    20
## 5 Revenge of the Sith   34
## 6 The Phantom Menace    34
## 7 Attack of the Clones  40
```

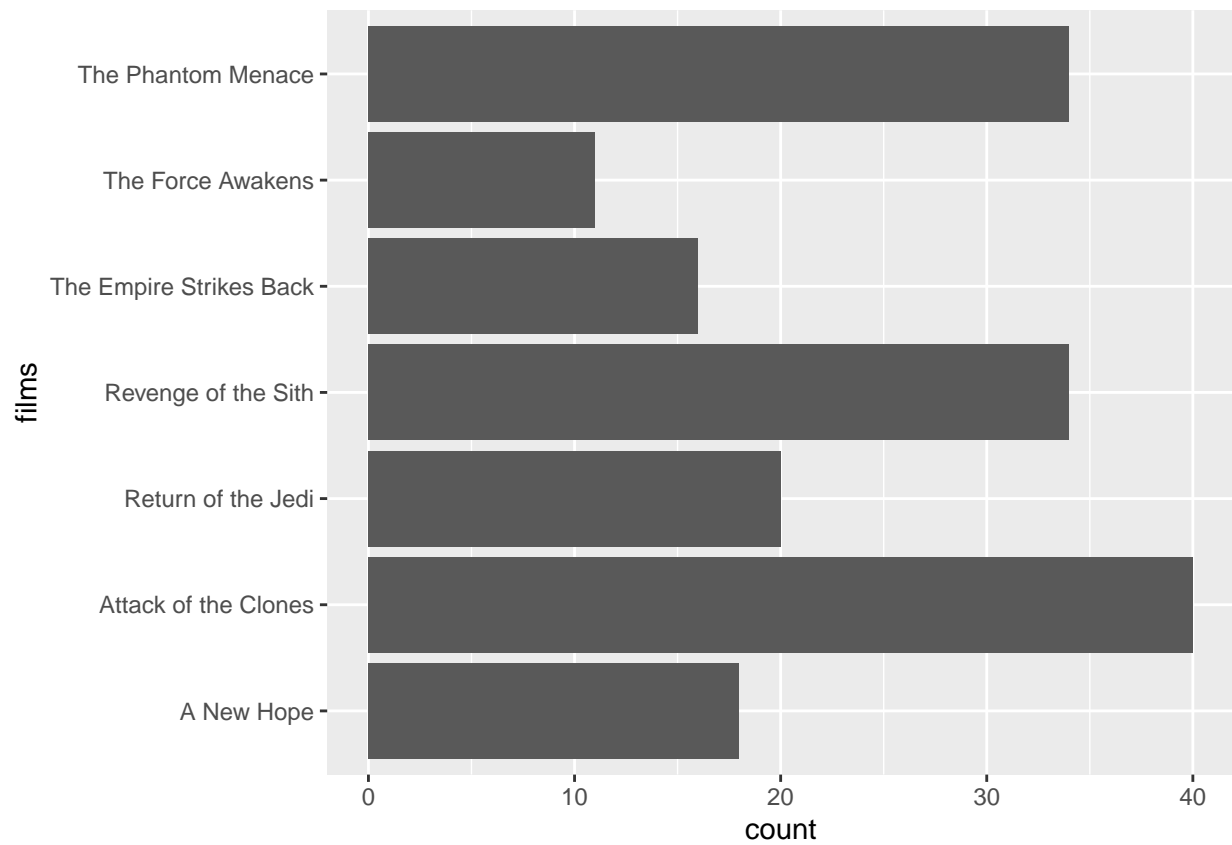
An other interesting function comes from forcats. In the previous example, we had a small number of a categories. However, quite often we will have A handful of common categories, and a whole bunch of other smaller groups. In such a case, we can use the forcats `fct_lump` to grab the most common categories, and lump the least most into a Other category.

```
starwars %>%
  filter(!is.na(homeworld)) %>%
  mutate(homeworld = fct_lump(homeworld, n = 3)) %>%
  count(homeworld) %>%
  arrange(n)
```

```
## # A tibble: 6 x 2
##   homeworld    n
##   <fct>      <int>
## 1 Alderaan     3
## 2 Coruscant    3
## 3 Kamino       3
## 4 Tatooine    10
## 5 Naboo       11
## 6 Other       47
```

Finally, we will demonstrate the `fct_infreq` function. In the first plot shown below, by default the plot is not ordered in any kind of way. However, by using `fct_infreq` in the second plot, we are able to reorder the values by their frequency in the data.

```
starwars %>%
  unnest(films) %>%
  ggplot(aes(films)) +
    geom_bar() +
    coord_flip()
```



```
starwars %>%  
  unnest(films) %>%  
  mutate(films = fct_infreq(films)) %>%  
  ggplot(aes(films)) +  
    geom_bar() +  
    coord_flip()
```

